



FUNDME

Product Operating System

Founder Brief, Product Architecture, and End-to-End Build Plan

Current baseline

Production is live at tryfundme.in. Accepted release: main at c363eb2. The next objective is to build the actual intelligence product - not continue polishing the landing page.

Prepared 21 June 2026

What this document does

This is the founder-facing control document for Fundme. It states what the product is, what exists now, what must be built, the order of execution, the operating model for autonomous agents, and the release gates that prevent context drift and repeated breakage.

The central shift

The landing page is no longer the project. It is one acquisition surface for a larger system: discover -> diagnose -> improve -> match -> draft -> track -> learn.

Contents

1. Executive decision summary
2. Current reality and gap map
3. Product thesis, user, and positioning
4. Competitive boundary and defensibility
5. Target product and technical architecture
6. Eight-week phase plan
7. Analytics, SEO/AEO, and distribution
8. Agent operating system and quality gates
9. Top risks and decisions
10. Immediate 14-day execution plan
11. Agent-document pack delivered with this brief

1. Executive decision summary

Fundme should become the application intelligence and execution layer for early-stage founders applying to accelerators, incubators, fellowships, startup competitions, selected grants, startup perks, and eventually investors.

The strongest initial product is not an auto-apply bot and not another opportunity directory. It is a reusable founder/startup dossier that supports evidence-backed diagnosis, explainable matching, program-specific drafting, and application tracking.

Product loop

Ingest -> evidence -> dossier -> diagnose -> improve -> match -> draft -> track -> outcome learning

Non-negotiable sequencing

1. Stabilize identity, environments, data ownership, monitoring, and analytics.
2. Build the founder/startup dossier and opportunity data model.
3. Build a versioned readiness assessment.
4. Build deterministic eligibility and weighted matching.
5. Build one credible application workspace.
6. Build tracking, reminders, outcomes, monetization, and referrals.
7. Only then expand into selective automation and network products.

What is deliberately frozen

- Broad visual redesign.
- Full autonomous browser submission.
- Broad grants bureaucracy and compliance workflows.
- Investor CRM and mass outreach.
- Community or social feed.
- Black-box startup scoring and unsupported AI claims.

2. Current reality and gap map

Area	Current state	Next required state
Production	Live and accepted at tryfundme.in	Keep stable while product core is built
Public discovery	Homepage plus public Explore/Search	Verified opportunity pages, richer dataset, SEO engine
Onboarding	Four-step early-access submission	Saved startup workspace and evidence-backed dossier
Authentication	Clerk integrated; production-instance transition still required	Production Clerk, identity sync, reliable staging
Database	Supabase submission persistence	Versioned domain schema, RLS,

		migrations, evidence model
Assessment	No canonical real engine	Versioned rubric plus grounded narrative
Matching	Directory and historical mock surfaces	Rules first, weighted score, explanation third
Drafting	Not a dependable product workflow	Question-level grounded drafts and versioning
Tracking	Historical routes may exist but require truth audit	Operational pipeline with deadlines and outcomes
SEO/AEO	Incomplete and unsystematic	Robots, sitemap, structured data, programmatic pages, content
Analytics	No canonical event architecture	End-to-end product and AI quality telemetry
Virality	Not built	Shareable report, referrals, outcome loop

Current highest-priority milestone

Platform foundation plus canonical startup dossier. Without these, every later AI feature becomes another disconnected demo.

3. Product thesis, user, and value

Primary user

- Pre-seed and seed founders applying to 5-30 programs in the next 30-90 days.
- Solo founders or small teams without a chief of staff.
- International and India-first founders navigating fragmented opportunity requirements.
- Founders with some material already, but low confidence in its quality and consistency.

Core jobs

1. Tell me how my startup and founder story reads to a decision-maker.
2. Show me what is weak or missing before I apply.
3. Tell me which opportunities are actually worth my time.
4. Reuse my verified context instead of making me rewrite everything.
5. Help me create a credible application specific to the program.
6. Keep deadlines, drafts, submissions, and outcomes organized.

Positioning

Primary

One startup dossier. Better applications everywhere.

Expanded

Fundme helps founders stress-test, improve, and deploy their startup story across the right opportunities.

North-star metric

Weekly founders completing an actionable funding workflow: a completed diagnosis plus at least one meaningful improvement, qualified match, or application draft.

4. Competitive boundary and defensibility

Product	Primary wedge	Strength	Fundme response
Fundraisly	Investor matching and outreach	Large investor dataset, warm paths, qualified meetings	Do not compete as a smaller outreach agent; own program application intelligence
VC Boom	Deck scoring, investor matching, outreach	Clear score-match-outreach funnel and fundraising expertise	Adopt clarity and evidence, but focus on accelerators, grants, perks, and application workflows
Fundme	Dossier, diagnosis, opportunity fit, application drafting, tracking	Reusable context plus structured program intelligence	Build evidence graph, outcome learning, India plus global depth

Defensibility

- Structured opportunity graph with requirements, questions, deadlines, benefits, sources, and freshness.
- Evidence graph containing verified founder/startup facts, source fragments, confidence, conflicts, and history.
- Program-question intelligence mapping the same evidence to different selection contexts.
- Outcome feedback from applications, interviews, rejections, acceptances, grants, perks, and funding.
- Workflow lock-in through reports, drafts, versions, deadlines, and application history.
- India plus global ecosystem depth that investor-only products do not prioritize.

5. Target product architecture

Product layers

Layer	What it contains
-------	------------------

Acquisition	Homepage, public opportunity pages, content, SEO, referrals
Source of truth	Founder profile, startup dossier, files, URLs, evidence, version history
Intelligence	Assessment, eligibility, matching, explanation, drafting
Execution	Application workspace, attachments, deadlines, tracker, outcomes
Commercial/network	Paid actions, credits, referrals, selective automation, ecosystem products

Technical direction

- Evolve the existing Next.js app as a modular monolith.
- Clerk for identity, Supabase for application data and storage.
- Typed route handlers/services with strict boundary validation.
- Queue long-running extraction, OCR, crawling, refresh, and bulk jobs only when needed.
- Provider-agnostic AI gateway with prompt/model/version metadata.
- Persistent staging environment to replace Preview sprawl.
- Structured logs, error monitoring, request IDs, cost and latency telemetry.

Core intelligence rule


Matching

Hard eligibility filters first -> weighted scoring second -> AI explanation third.

Drafting

Verified facts + program requirements -> tailored draft. Missing evidence creates a question or warning, never an invented claim.

6. Eight-week phase plan



Phase	Outcome	Exit gate
0 - Truth lock	Install docs, audit repo, lock backlog	A new agent can start without chat history
1 - Platform	Production Clerk, staging, identity, RLS, monitoring, analytics	Auth and environments are dependable
2A - Data/discovery	Opportunity schema, admin, 100 quality records target, SEO basics	Useful indexable opportunity pages
2B - Dossier	Uploads, parsing, evidence facts, editable versioned dossier	Founder has a reusable source of truth
3 - Assessment	Versioned rubric and evidence-backed report	Representative QA set passes factuality
4 - Matching	Eligibility, weighted scoring, explanation	Scores are reproducible and ineligible records excluded
5 - Drafting	Program questions, grounded answers, versions, export	One real application package completed
6 - Tracker	Statuses, deadlines, reminders, outcomes	User knows the next action from dashboard
7 - Commercial	Entitlements, payments, referrals, lifecycle	Product can monetize demonstrated intent

A single serial agent should plan for 10-12 weeks. Dates are aggressive targets and never override exit gates.

7. SEO, AEO, analytics, and distribution

Organic discovery

- Canonical public page for every verified opportunity.
- Category, geography, stage, and sector hubs only when they provide unique value.
- Robots.txt, sitemap, canonical metadata, structured data, Search Console, and internal links.
- Question-led guides: what programs look for, deadline cycles, application breakdowns, and comparisons.
- Experimental llms.txt for machine orientation, not as a ranking promise.
- Partnership backlinks from universities, incubators, founder communities, perk providers, and ecosystem newsletters.

Analytics

Funnel	Critical events
Acquisition	landing_viewed, program_page_viewed, primary_cta_clicked
Dossier	assessment_started, source_added, dossier_generated, dossier_fact_corrected
Assessment	assessment_submitted, assessment_generated, report_viewed, improvement_action_clicked
Matching	match_run_completed, match_viewed, match_shortlisted, match_dismissed
Drafting	draft_started, draft_generated, draft_edited, draft_exported
Outcomes	application_marked_submitted, status_changed, outcome_recorded
Commercial/viral	paywall_viewed, subscription_started, report_shared, referral_activated

AI quality must be measured alongside conversion: unsupported claim rate, fact-correction rate, parse failures, ineligible match exposure, draft edit distance, cost, and latency.

8. Autonomous-agent operating system

Source-of-truth hierarchy

1. Current repository and production behavior.
2. PROJECT_STATE.md.
3. Canonical specifications.
4. Recorded decisions.
5. Current phase contract.
6. Historical prompts and archived reports.

How agent autonomy works

Agents may choose implementation details, refactor within scope, add tests, and repair directly related defects. They must escalate only product-direction changes, pricing, destructive migrations, privacy/security tradeoffs, and production releases outside policy.

Quality rule

Evidence standard
Code inspection is not browser evidence. A PASS requires an executed browser state, resulting URL, network response, safe database proof, test output, or recorded walkthrough.

Model roles

Agent	Role
Codex	Primary builder: implementation, types, migrations, tests, handoff
Antigravity	Browser/visual QA, environment verification, release evidence; builder only when assigned
ChatGPT	Product architect, plan critic, document maintainer, phase-contract author, decision orchestrator

9. Top risks and control decisions

Risk	Control
Context drift across agents	Canonical docs, project-state file, phase contracts, update-on-merge rule
Landing page consumes all effort	Functional-core roadmap and redesign freeze
AI invents founder facts	Evidence model, strict schemas, user verification, unsupported-claim eval
Stale opportunity data	Sources, last_verified_at, admin queue, stale labels, user corrections
Matching becomes AI vibes	Eligibility rules and versioned weighted scoring
Deployment/environment sprawl	Persistent staging, controlled branch policy, environment parity checks
Security key exposure	Server-only secrets, no logs/prompts, key migration and rotation
Broad scope kills velocity	Explicit non-goals and evidence-gated later phases
Generic competitor clone	Own application intelligence and program opportunity graph

10. Immediate 14-day execution plan

Days 1-3: truth and foundation audit

- Install the full Markdown pack in the repository.
- Audit current app routes, mocks, APIs, database schema, environments, and AI usage.
- Update PROJECT_STATE.md from repository evidence.
- Prepare the persistent staging design and the Phase 1 implementation diff.

Days 4-8: production identity and operations

- Complete Clerk production-instance migration and branding.
- Implement Clerk-to-Supabase identity mapping and authorization checks.
- Audit RLS and elevated-key use.
- Add error monitoring, structured API errors, and environment-name preflight.

Days 9-14: analytics and dossier foundation

- Implement core acquisition/onboarding analytics.
- Create versioned startup, founder, source, and evidence schemas.
- Build authenticated startup workspace shell and manual dossier editor.
- Prepare ingestion fixtures and Phase 2 contracts.

What not to do in these 14 days

Do not redesign the homepage, add browser automation, launch broad grants, or build payment before the identity and dossier foundation is accepted.

11. Agent document pack

The accompanying Markdown package contains the detailed, machine-readable source of truth for agents. It includes the PRD, architecture, AI and matching spec, SEO/AEO, analytics, design, security, delivery roadmap, quality playbook, agent adapters, decisions, templates, and the first phase contract.

File group	Use
Core context	README, PROJECT_STATE, FUNDME_PROJECT_OS, PRODUCT_CHARTER
Product/UX	MASTER_PRD, USER_FLOWS_UX, DESIGN_CONTENT_SYSTEM
Engineering	TECHNICAL_ARCHITECTURE, DATA_AI_MATCHING_SPEC, SECURITY_PRIVACY_OPERATIONS
Growth	OPPORTUNITY_DATA_SEO_AEO, ANALYTICS_GROWTH_VIRALITY

Execution	DELIVERY_ROADMAP, QUALITY_RELEASE_PLAYBOOK, PHASE_01_BUILD_CONTRACT
Agents	AGENTS, CODEX, ANTIGRAVITY, CHATGPT_ORCHESTRATOR
Governance	DECISIONS, SOURCES_AND_ASSUMPTIONS, templates

References and current technical notes

- Fundraisly: <https://fundraisly.com/>
- VC Boom: <https://www.vcboom.com/>
- Next.js proxy convention: <https://nextjs.org/docs/app/api-reference/file-conventions/proxy>
- Clerk production deployments: <https://clerk.com/docs/deployments/overview>
- Clerk environments: <https://clerk.com/docs/deployments/environments>
- Supabase API keys: <https://supabase.com/docs/guides/getting-started/api-keys>
- Vercel environment variables: <https://vercel.com/docs/environment-variables>
- Google AI features and websites: <https://developers.google.com/search/docs/appearance/ai-features>
- Google robots.txt: <https://developers.google.com/search/docs/crawling-indexing/robots/create-robots-txt>
- Google structured data: <https://developers.google.com/search/docs/appearance/structured-data/search-gallery>
- llms.txt proposal: <https://llmstxt.org/>

This brief also distills the Fundme Master PRD, source-of-truth extraction, clean MVP handoff, VUX plans, production acceptance history, and operational lessons supplied in the project conversation and files.

Final directive

From this point forward, agents build against the repository-installed operating system and current phase contract. Chat history is supporting context, not the source of truth.